

Analysis of Venues Around MRT Stations in Singapore Using Foursquare Location Data

By: Choy Siew Fong

(Applied Data Science Capstone Project Submitted as part of the Requirements for IBM Professional Certificate in Data Science)

23 March 2020

1. Introduction

Singapore's rail network system consists of the Mass Rapid Transit (MRT), which is a heavy rail rapid transit system and the subsidiary light rail transit (LRT) system which serves the non-mature estates of Sengkang, Punggol and Bukit Panjang and function as a feeder network to the MRT. Figure 1 shows an example of an MRT train while Figure 2 shows a typical LRT train.



Figure 1: Photograph of an MRT train at Mandai Depot (Source: Mass Rapid Transit (Singapore). 18 March 2020. In *Wikipedia*. Retrieved March 19, 2020, [https://en.wikipedia.org/wiki/Mass_Rapid_Transit_\(Singapore\)](https://en.wikipedia.org/wiki/Mass_Rapid_Transit_(Singapore)))



Figure 2: An LRT train (Source: Light Rail Transit (Singapore). 18 March 2020. In *Wikipedia*. Retrieved March 19, 2020, [https://en.wikipedia.org/wiki/Light_Rail_Transit_\(Singapore\)](https://en.wikipedia.org/wiki/Light_Rail_Transit_(Singapore)))

Due to the astronomical prices of cars in Singapore, most of the island's population gets around by public transport. In 2019, the daily ridership of the MRT network was 3.4 million while the annual ridership was 1.2 billion. Since the MRT and LRT is such an integral part of people's daily lives in Singapore, I was interested to find out what kind of venues (such as restaurants, cafes, or supermarkets et cetera) existed around the MRT and LRT stations and if they have any trends or commonalities which may be analysed using clustering methods such as KMeans. This information may be obtained by combining a dataset of the stations' latitude and longitude coordinates with Foursquare location data.

1.1 Potential Stakeholders

This information about the type and number of venues available around each MRT/LRT station would be of interest to various stakeholders. Firstly, passengers or commuters taking the train would be keen to know what amenities are around each MRT or LRT station so that they can make informed decisions during their commute, such as which station to drop off if they wish to buy a loaf of bread from the supermarket for tomorrow's breakfast or which station would have an ATM from their bank.

Application developers may also have an interest in this information as it will enable them to develop mobile applications which can locate a certain venue (e.g. ATMs or cafes) which would be very useful to potential users.

Potential business owners who are contemplating starting a business around MRT stations due to their high footfall would also be interest in this information since it would allow them to know the number and type of competitors already existing in the area. For instance, a businessman interested in opening a Japanese restaurant in the Jurong East MRT area can use this information to find out how many *other* Japanese restaurants already exist in the area, and what are their specialities, if any. This information would allow prospective business owners to make informed decisions on which niche area they wish to target or to decide to move to a different location altogether.

Similarly, current business owners would be interested in knowing which venues are trending in the locale of their business. This would highlight potential competitors or trends which would enable them to adjust their business model and menu accordingly.

2. Description of the Dataset

The dataset 'Singapore Train Station Coordinates' was obtained from Kaggle (see <https://www.kaggle.com/yxlee245/singapore-train-station-coordinates>). It is in comma separated values (.csv) format and consists of four columns listing the station name, type (i.e. whether it is an MRT or an LRT station) and positional coordinates (latitude and longitude) of each MRT and LRT station in Singapore at the time of its upload, which was around eight months ago. Figure 3 lists the top 5 rows of the dataset.

	station_name	type	lat	lng
1	Jurong East	MRT	1.333207	103.742308
2	Bukit Batok	MRT	1.349069	103.749596
3	Bukit Gombak	MRT	1.359043	103.751863
4	Choa Chu Kang	MRT	1.385417	103.744316

Figure 3: Screenshot from Excel showing the top 5 rows of the dataset

Preliminary inspection of the dataset found no missing or null values. However, it was later discovered during Folium map plotting that the coordinates for Admiralty Station were identical to Woodlands Station, probably due to a typographical error by the author of the dataset. To resolve this issue, I searched online

for the coordinates of Admiralty station and replaced the erroneous longitude and latitude values for Admiralty station in the dataset with those from this site* (latitude = 1.44069°, longitude = 103.8009°).

*(<https://www.findlatitudeandlongitude.com/l/Admiralty+mrt/1214628/>).

The Foursquare information about each station would be obtained by defining a search URL (limiting the number of venues returned for each station to 100 within a radius of 500 m from each station) and sending the **GET** request to Foursquare API.

3. Methodology section

3.1 Exploratory Data Analysis

Using the **pd.read_csv** function, the dataset was first read into a Pandas dataframe, named **df**, as follows. It was noted that the column names were 'station_name', 'type', 'lat' and 'lng' (for latitude and longitude) as shown in Figure 4 below.

In [2]:

read data file
df = pd.read_csv('mrt_lrt_data2.csv')
df.head()

Out[2]:

	station_name	type	lat	lng
0	Jurong East	MRT	1.333207	103.742308
1	Bukit Batok	MRT	1.349069	103.749596
2	Bukit Gombak	MRT	1.359043	103.751863
3	Choa Chu Kang	MRT	1.385417	103.744316
4	Yew Tee	MRT	1.397383	103.747523

Figure 4: Top 5 rows of the dataset when read into a dataframe **df**

To understand more about the dataset, the **dtypes** command was used to find out the data type of each column. This enables us to determine which functions can be applied and to do any conversions to other data types, if needed. As shown in Figure 5, the 'station_name' and 'type' fields were found to be of type *Object* while the 'lat' and 'lng' fields were of type *Float*.

<pre># Basic type information about data df.dtypes</pre>	
station_name	object
type	object
lat	float64
lng	float64
dtype:	object

Figure 5: Basic data type information of the dataset obtained with **dtypes** command

The data was then grouped by type of stations using the **groupby** function and the number of each station type obtained with the **count** function. There are 119 MRT stations and 38 LRT stations in the dataset.

```
#check number of LRT and MRT stations
df_group = df.groupby('type').count()
df_group
```

	station_name	lat	lng
type			
LRT	38	38	38
MRT	119	119	119

Figure 6: Output showing the number of each type of stations

It was decided at this point to confine the subsequent analysis to MRT stations only as these were the main stations while LRT stations are in the suburbs and may have few interesting venues around them if any. Moreover, the free Foursquare developer sandbox account allows only 950 regular calls a day. Using the full dataset may lead to issues with hitting the call limit later, given that each station is expected to have a substantial number of venues.

A new dataframe containing only the MRT stations was then defined as follows.

```
#Get MRT data into new dataframe - confine analysis to MRT stations
df_MRT = df.loc[df['type'] == 'MRT']
df_MRT.head()
```

Figure 7: Defining new dataframe df_MRT containing only MRT station data.

The shape of the new dataframe was checked using the **shape** attribute as below.

```
In [6]: # check size of MRT data
df_MRT.shape

Out[6]: (119, 4)
```

Figure 8: Checking shape of the resulting dataframe

This output confirms that the new **df_MRT** dataframe did contain information from all 119 MRT stations that were in the original dataset. The data was then visualised as a map using **Folium**.

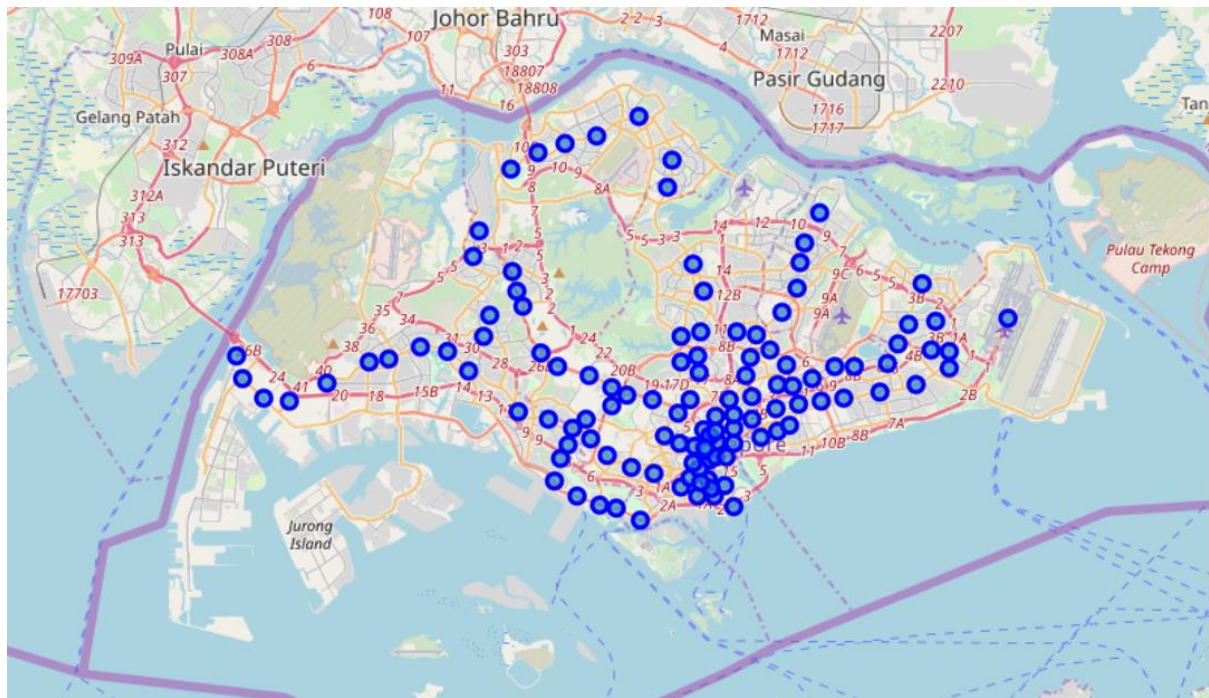


Figure 9: Folium map showing the locations of 119 MRT stations

3.2 Exploring the first MRT station

It is time to start using Foursquare API to get location data about the stations. I began the analysis by defining my Foursquare credentials (client ID, client secret and version). The first station in the dataframe was then identified using `.loc` command, with zero as the first row index.

```
#first station location
df_MRT.loc[0, 'station_name']

'Jurong East'
```

Figure 10: Finding the first station name from df_MRT

The latitude and longitude of Jurong East station was defined and a **GET** URL request sent to get information of 100 venues located within 500 m from the station.

```
In [12]: results = requests.get(url).json() #get results
results

Out[12]: {'meta': {'code': 200, 'requestId': '5e74dc66660a9f001ea09003'},
  'response': {'suggestedFilters': {'header': 'Tap to show:',
    'filters': [{'name': 'Open now', 'key': 'openNow'}]},
  'headerLocation': 'Jurong East',
  'headerFullLocation': 'Jurong East, Singapore',
  'headerLocationGranularity': 'neighborhood',
  'totalResults': 76,
  'suggestedBounds': {'ne': {'lat': 1.3377070045000046,
    'lng': 103.74680081871648},
    'sw': {'lat': 1.3287069954999955, 'lng': 103.73781518128351}},
  'groups': [{'type': 'Recommended Places',
    'name': 'recommended',
    'items': [{'reasons': {'count': 0,
```

Figure 11: Json file output by the GET request

A get category type function was defined and the json output normalised. The resulting dataframe was then filtered to display only location related data such as venue name, venue category and venue coordinates (as latitude and longitude) (Figure 12). The output was then checked by calling the **head** function (Figure 13).

```
In [13]: #define get_category_type function
def get_category_type(row):
    try:
        categories_list = row['categories']
    except:
        categories_list = row['venue.categories']
    if len(categories_list) == 0:
        return None
    else:
        return categories_list[0]['name']

In [14]: venues = results['response']['groups'][0]['items']
nearby_venues = json_normalize(venues)

In [15]: # filter columns
filtered_columns = ['venue.name', 'venue.categories', 'venue.location.lat', 'venue.location.lng']
nearby_venues = nearby_venues.loc[:, filtered_columns]
# filter the category for each row
nearby_venues['venue.categories'] = nearby_venues.apply(get_category_type, axis=1)
# clean columns
nearby_venues.columns = [col.split(".")[1] for col in nearby_venues.columns]
nearby_venues.head()
```

Figure 12: Defining the get category type function and filtering and cleaning the columns

```
Out[15]:
```

	name	categories	lat	lng
0	UNIQLO	Clothing Store	1.333175	103.743160
1	MUJI 無印良品	Furniture / Home Store	1.333187	103.743064
2	Song Fa Bak Kut Teh 松發肉骨茶	Chinese Restaurant	1.333394	103.743420
3	Johan Paris	Bakery	1.334083	103.742384
4	The Rink	Skating Rink	1.333424	103.740345

Figure 13: First 5 rows of nearby_venues dataframe

The venues found around Jurong East station was visualised using **Folium.Map** function. A red marker was added to indicate the location of the MRT station while blue markers were used to denote the venues.

```
In [16]: #Map for regions
venues_map = folium.Map(location=[station_latitude, station_longitude], zoom_start=20)
# add a red circle marker to represent the Jurong East station
folium.features.CircleMarker(
    [station_latitude, station_longitude],
    radius=10,
    color='red',
    popup='Jurong East',
    fill = True,
    fill_color = 'red',
    fill_opacity = 0.6
).add_to(venues_map)
# add all venues as blue circle markers
for lat, lng, label in zip(nearby_venues.lat, nearby_venues.lng, nearby_venues.categories):
    folium.features.CircleMarker(
        [lat, lng],
        radius=5,
        color='blue',
        popup=label,
        fill = True,
        fill_color='blue',
        fill_opacity=0.6
    ).add_to(venues_map)

venues_map
```

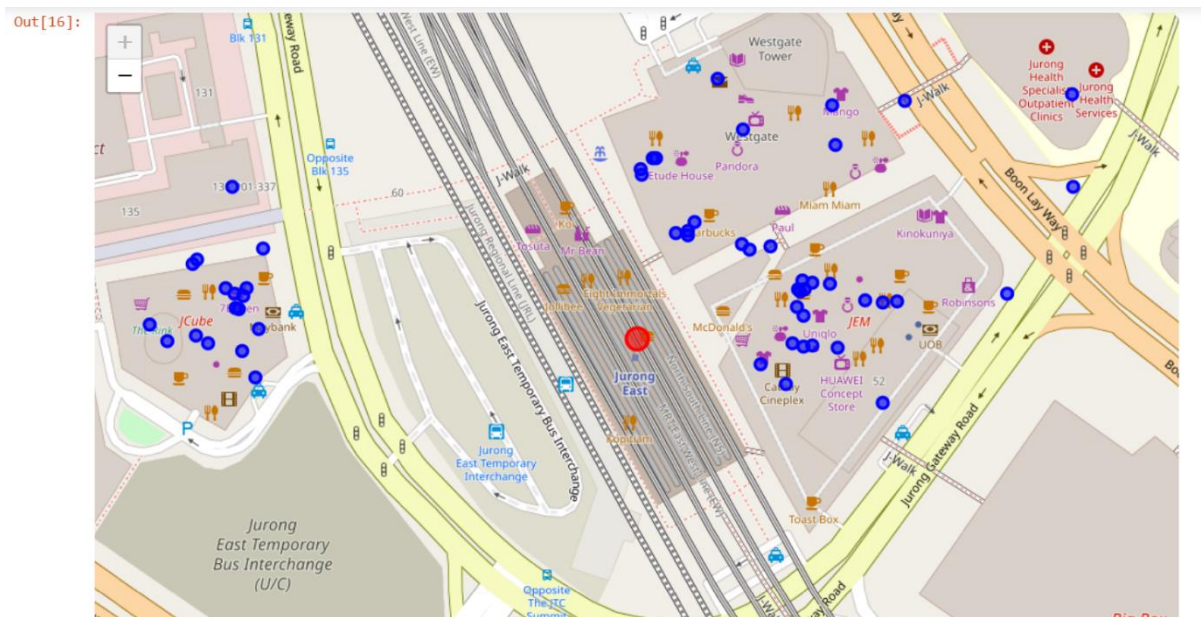



Figure 14: Map of venues found from Foursquare located within 500 m around the first MRT station, Jurong East.

Using the **shape** function on nearby_venues dataframe, it was found that a total of 76 venues were returned around Jurong East station by Foursquare.

```
In [17]: print('{} venues were returned by Foursquare.'.format(nearby_venues.shape[0])) # number of venues returned
76 venues were returned by Foursquare.
```

Figure 15: Total number of venues returned by Foursquare

3.3 Expanding the analysis to the full dataset

Now the same analysis can be performed on all the MRT stations contained in df_MRT. First, a getNearbyVenues function was defined. It was then applied to df_MRT and the results were output to a new dataframe called station_venues.

```
station_venues = getNearbyVenues(names=df_MRT['station_name'],
                                  latitudes=df_MRT['lat'],
                                  longitudes=df_MRT['lng'])
```

Figure 16: Applying the defined function getNearbyVenues to df_MRT

Once again, the shape function was used to check the size of the resultant dataframe. It was found that the station_venues dataframe consisted of 4723 rows and 7 columns.

```
In [20]: #check size of station_venues
print(station_venues.shape)
station_venues.sample(10)

(4723, 7)
```

Figure 17: Shape of station_venues dataframe

To check the output of the dataframe station_venues, the **sample** function was used to return a random sample of 10 stations' data. Using the **head** function will return only the first one or two station's data and is not as representative of the full dataframe as the sample function.

Out[20]:

	Station	Station_Latitude	Station_Longitude	Venue	Venue_Latitude	Venue_Longitude	Venue_Category
1051	Marina Bay	1.276481	103.854598	Marina Bay Financial Centre (MBFC) Tower 3	1.279109	103.854420	Building
2114	Tampines	1.354467	103.943325	J's Wok & Grill	1.352120	103.941506	Mediterranean Restaurant
440	Braddell	1.340550	103.847098	Chey Sua Carrot Cake	1.338012	103.844661	Chinese Breakfast Place
395	Bishan	1.350920	103.848206	Gourmet Food Court	1.350661	103.850668	Food Court
2670	Farrer Park	1.312679	103.854872	Machan's Kitchen	1.310036	103.851390	Restaurant
4053	Rochor	1.303601	103.852581	DECK	1.301771	103.851645	Art Gallery
3298	Stadium	1.302847	103.875417	Popeyes Louisiana Kitchen	1.303195	103.873082	Fried Chicken Joint
4399	Bencoolen	1.298477	103.849984	Toast Box	1.298130	103.849891	Café
1214	Lakeside	1.344264	103.720797	7-Eleven	1.346075	103.718105	Grocery Store
1271	Clementi	1.314925	103.765341	Prata Alley	1.312084	103.765132	Asian Restaurant

Figure 18: Sample containing 10 rows data from station_venues dataframe.

3.4 Analysing the station_venues dataframe

The station_venues dataframe would be the main dataframe being analysed from this point onwards as it contains the combined data of both MRT station names and location, as well as nearby venues name, category and location.

First, preliminary analysis was performed using the **unique** function to find out the number of unique venues and venue category being returned by Foursquare.

```
In [21]: #find unique categories number
print('There are {} unique venue categories.'.format(len(station_venues['Venue_Category'].unique())))

There are 322 unique venue categories.

In [22]: #find unique venue number
print('There are {} uniques venues.'.format(len(station_venues['Venue'].unique())))#number of unique venues

There are 3223 uniques venues.
```

Figure 19: Number of unique venue categories and unique venues

A total of 322 *unique venue categories* and 3223 *unique venues* were found around 119 MRT stations in the dataset. The total number of venues returned for each MRT station was determined using the **groupby** and **count** functions as follows.

```
station_venues.groupby('Station').count()
```

Figure 20: Using groupby and count functions to find the total number of venues for each MRT station

Out[23]:

	Station_Latitude	Station_Longitude	Venue	Venue_Latitude	Venue_Longitude	Venue_Category
Station						
Admiralty	9	9	9	9	9	9
Aljunied	50	50	50	50	50	50
Ang Mo Kio	40	40	40	40	40	40
Bartley	10	10	10	10	10	10
Bayfront	50	50	50	50	50	50
Beauty World	78	78	78	78	78	78
Bedok	57	57	57	57	57	57
Bedok North	17	17	17	17	17	17
Bedok Reservoir	8	8	8	8	8	8
Bencoolen	100	100	100	100	100	100

Figure 21: Screenshot showing the first 10 rows of the station_venues dataframe

To perform further analysis on the station_venues dataframe, it is necessary to convert the categorical values for Venue_Category to numerical values (i.e. one hot vectors). This conversion (one hot encoding) was done using the **get_dummies** function (see Jupyter notebook for details). The rows in the resulting dataframe were then grouped by station and the mean of the frequency of occurrence of each category. The resultant dataframe, called station_grouped, had 119 rows and 323 columns.

	Station	ATM	Accessories Store	Airport	Airport Lounge	American Restaurant	Arcade	Art Gallery	Art Museum	Arts & Crafts Store	...	Water Park	Waterfall	Waterfront
0	Admiralty	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000
1	Aljunied	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000
2	Ang Mo Kio	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000
3	Bartley	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000
4	Bayfront	0.000000	0.020000	0.000000	0.000000	0.000000	0.000000	0.020000	0.00	0.000000	...	0.000000	0.000000	0.040000
5	Beauty World	0.012821	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000
6	Bedok	0.000000	0.000000	0.000000	0.000000	0.017544	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000
7	Bedok North	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00	0.000000	...	0.000000	0.000000	0.000000

Figure 22: Screenshot showing station_grouped dataframe

The top 5 venue category around each station were printed to get an idea of which venues were common as shown below.

```
In [28]: #print each station along with the top 5 most common venues
num_top_venues = 5
for stn in station_grouped['Station']:
    print("----"+stn+"----")
    temp = station_grouped[station_grouped['Station'] == stn].T.reset_index()
    temp.columns = ['venue', 'freq']
    temp = temp.iloc[1:]
    temp['freq'] = temp['freq'].astype(float)
    temp = temp.round({'freq': 2})
    print(temp.sort_values('freq', ascending=False).reset_index(drop=True).head(num_top_venues))
    print('\n')
```

```
----Admiralty----
      venue  freq
0  Supermarket  0.22
1   Food Court  0.11
2  Bus Station  0.11
3         Spa  0.11
4         Park  0.11

----Aljunied----
      venue  freq
0  Chinese Restaurant  0.12
1   Noodle House  0.10
2    Coffee Shop  0.08
3  Asian Restaurant  0.08
4    Food Court  0.06
```

Figure 23: Printing the top 5 venue categories for each station.

The venue categories were then sorted in descending order and a new dataframe, station_venues_sorted, containing the top ten venue categories for each MRT station created as shown below.

Out[30]:

	Station	1st Most Common Venue	2nd Most Common Venue	3rd Most Common Venue	4th Most Common Venue	5th Most Common Venue	6th Most Common Venue	7th Most Common Venue	8th Most Common Venue	9th Most Common Venue	10th Most Common Venue
0	Admiralty	Supermarket	Park	Spa	Convenience Store	Plaza	Bus Station	Food Court	Coffee Shop	Food & Drink Shop	Field
1	Aljunied	Chinese Restaurant	Noodle House	Asian Restaurant	Coffee Shop	Vegetarian / Vegan Restaurant	Food Court	Dim Sum Restaurant	Breakfast Spot	Bus Station	Seafood Restaurant
2	Ang Mo Kio	Coffee Shop	Dessert Shop	Food Court	Supermarket	Bubble Tea Shop	Japanese Restaurant	Multiplex	Ramen Restaurant	Miscellaneous Shop	Fast Food Restaurant
3	Bartley	Bus Station	Noodle House	Seafood Restaurant	Metro Station	Café	Pet Store	Yunnan Restaurant	Food Court	Fish & Chips Shop	Flea Market
4	Bayfront	Hotel	Boutique	Japanese Restaurant	Waterfront	Bar	Theater	Italian Restaurant	Bridge	Tea Room	Lounge

Figure 24: First 5 rows of the station_venues_sorted dataframe showing top 10 venue category for each station

3.5 KMeans Clustering

The station_grouped dataframe was then analysed using KMeans clustering with a cluster size of 5 after dropping the first column 'Station' which had a non-numerical value. The resulting dataframe was then merged with the original df_MRT to include the latitude and longitude data. Figure 25 below shows the relevant part of the code.

```
In [31]: # set number of clusters
kclusters = 5
station_grouped_clustering = station_grouped.drop('Station', 1)
# run k-means clustering
kmeans = KMeans(n_clusters=kclusters, random_state=0).fit(station_grouped_clustering)
# check cluster labels generated for each row in the dataframe
kmeans.labels_[0:10]

Out[31]: array([2, 3, 2, 4, 0, 3, 0, 2, 2, 0])

In [32]: # add clustering labels
station_venues_sorted.insert(0, 'Cluster Labels', kmeans.labels_)
station_merged = df_MRT
# merge station_grouped with station_data to add Latitude/Longitude for each station
station_merged = station_merged.join(station_venues_sorted.set_index('Station'), on='station_name')
station_merged.head() # check dataframe
```

Figure 25: Clustering using the station_grouped dataframe

The clusters were then visualised in a Folium map as shown below.

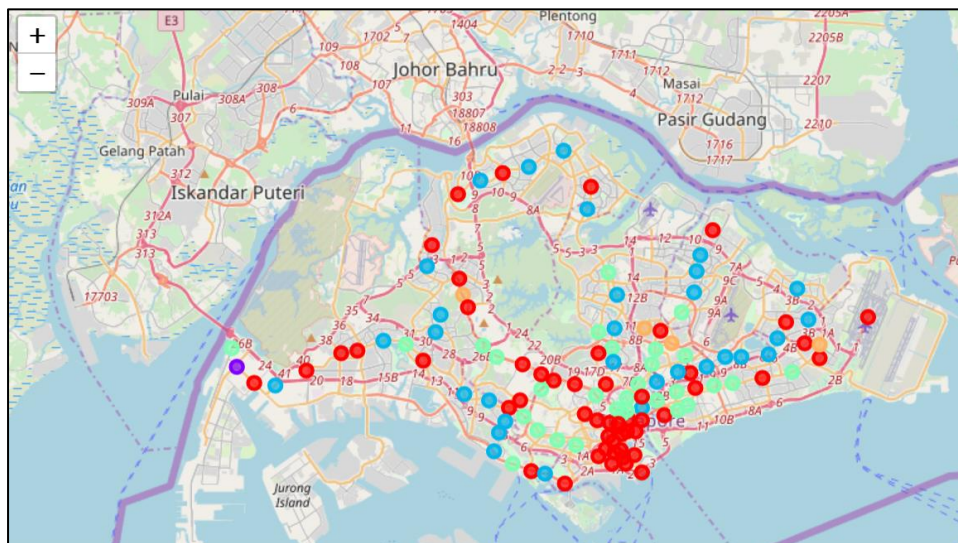


Figure 26: Folium map showing the 5 clusters

3.6 Determining the Station with the Highest Venue Category

Since each station can have a variable number of venue categories and a certain number of each category, I would like to find out which station has the maximum number of a venue category, and what category it is. This was done using the **groupby** and **size** functions.

```
In [37]: #group by station and find number of venue category of each type
stn_grp = station_venues.groupby(['Station', 'Venue_Category']).size()
stn_grp
```

```
Out[37]: Station  Venue_Category
Admiralty  Bus Station          1
           Coffee Shop          1
           Convenience Store    1
           Food Court           1
           Park                 1
           Plaza                1
           Spa                  1
           Supermarket          2
Aljunied   Asian Restaurant     4
           BBQ Joint            2
           Badminton Court      1
           Boarding House       1
           Breakfast Spot       2
           Bus Station          2
           Café                 1
           Chinese Restaurant   6
           Coffee Shop          4
           Convenience Store    1
           Dim Sum Restaurant    2
           Farmers Market       1
```

Figure 27: Grouping stations by station and Venue Category with resultant output

The station with the maximum number of venue category was then listed using the **idxmax** function.

```
In [38]: #find station with max number of venue category
stn_grp.idxmax()

Out[38]: ('Jalan Besar', 'Indian Restaurant')
```

Figure 28: Station with maximum number of a venue category

Surprisingly, Jalan Besar returned the greatest number of venue category (Indian restaurant). The Indian restaurants in Jalan Besar were called from the station_venues dataframe. A total of 21 indian restaurants were returned (see Figure 29).

Out[39]:

	Station	Station_Latitude	Station_Longitude	Venue	Venue_Latitude	Venue_Longitude	Venue_Category
4493	Jalan Besar	1.305551	103.855443	Bismillah Biryani	1.304956	103.853602	Indian Restaurant
4499	Jalan Besar	1.305551	103.855443	Murugan Idli Shop	1.308842	103.856380	Indian Restaurant
4506	Jalan Besar	1.305551	103.855443	Azmi Restaurant	1.308256	103.853075	Indian Restaurant
4507	Jalan Besar	1.305551	103.855443	Sakunthala's Restaurant	1.306000	103.852169	Indian Restaurant
4513	Jalan Besar	1.305551	103.855443	Komala Vilas (Buffalo Rd)	1.306308	103.851158	Indian Restaurant
4516	Jalan Besar	1.305551	103.855443	Khansama Tandoori Restaurant	1.308251	103.853122	Indian Restaurant
4526	Jalan Besar	1.305551	103.855443	Kailash Parbat	1.308039	103.852660	Indian Restaurant
4527	Jalan Besar	1.305551	103.855443	Lagnaa Barefoot Dining	1.306472	103.852298	Indian Restaurant
4535	Jalan Besar	1.305551	103.855443	Veeras Curry Restaurant @ Hindoo Rd	1.308650	103.853515	Indian Restaurant
4539	Jalan Besar	1.305551	103.855443	Sakunthala's Restaurant	1.309475	103.855717	Indian Restaurant
4546	Jalan Besar	1.305551	103.855443	Premaas Cuisine	1.305094	103.851980	Indian Restaurant
4556	Jalan Besar	1.305551	103.855443	Taste Of India	1.307900	103.852089	Indian Restaurant
4559	Jalan Besar	1.305551	103.855443	The Banana Leaf Apolo	1.305348	103.851698	Indian Restaurant
4561	Jalan Besar	1.305551	103.855443	Islamic Restaurant	1.303075	103.859172	Indian Restaurant
4563	Jalan Besar	1.305551	103.855443	Kebabs 'n Curries	1.309157	103.856456	Indian Restaurant
4566	Jalan Besar	1.305551	103.855443	Allaaddin's Briyani	1.305799	103.851180	Indian Restaurant
4567	Jalan Besar	1.305551	103.855443	Ma Raj Restaurant	1.309817	103.855970	Indian Restaurant
4568	Jalan Besar	1.305551	103.855443	Saravanaa Bhavan	1.309272	103.855971	Indian Restaurant
4570	Jalan Besar	1.305551	103.855443	Copper Chimney	1.309712	103.855447	Indian Restaurant
4575	Jalan Besar	1.305551	103.855443	Ananda Bhavan Restaurant	1.309665	103.855614	Indian Restaurant
4580	Jalan Besar	1.305551	103.855443	Yakader	1.305794	103.851108	Indian Restaurant

Figure 29: Total number of Indian restaurants around Jalan Besar station.

3.7 Top 10 Venue Category for all Stations

To determine the top ten most common venue category for all MRT stations, the `stn_grp` object returned in Figure 27 was converted to a dataframe, which was grouped by Venue Category and the counts for each venue category summed up. The resultant output is as follows:

Out[41]:

Venue_Category	
ATM	2
Accessories Store	4
Airport	2
Airport Lounge	5
American Restaurant	23
Arcade	4
Art Gallery	22
Art Museum	2
Arts & Crafts Store	11
Asian Restaurant	126
Athletics & Sports	3
Australian Restaurant	5
BBQ Joint	44
Baby Store	1
Badminton Court	1
Bagel Shop	2
Bakery	124
Bank	2
Bar	43
Basketball Court	6
Bay	1
Bed & Breakfast	5
Beer Bar	16
Beer Garden	8
Beer Store	1
Betting Shop	3
Bike Rental / Bike Share	1
Bike Shop	1
Bistro	20
Board Shop	1
...	

Figure 30: Count for each venue category for all MRT stations

This list was then sorted in descending order and the top ten venue categories displayed using the **head(10)** command.

```
In [42]: Alist.sort_values(ascending = False, inplace = True) #sort venue category in descending order
Top_10 = Alist.head(10) #find top 10 venue category
Top_10

Out[42]: Venue_Category
Coffee Shop          251
Chinese Restaurant   239
Café                 229
Food Court           177
Japanese Restaurant  172
Hotel                132
Asian Restaurant     126
Bakery               124
Indian Restaurant    115
Fast Food Restaurant  96
Name: Counts, dtype: int64
```

Figure 31: Top 10 list of venue category

To visualise the results, matplotlib bar chart was used to plot the top 10 categories.

```
In [45]: # plot Bar graph
import matplotlib as mpl
import matplotlib.pyplot as plt

Top_10.plot(kind = 'bar')

plt.title('Top Ten Venue Category in MRT and LRT Stations')
plt.xlabel('Venue Category')
plt.ylabel('Number')
plt.show()
```




Figure 32: Bar chart showing top 10 venue categories for all MRT stations

3.8 Determining the Station with the Maximum Number of a Venue

A similar analysis was done to determine the station with the highest number of one particular venue. It turned out that Downtown station had the maximum number of Starbucks as shown below.


```
In [47]: stn_grp2.idxmax() #Find venue with max number and the corresponding station
Out[47]: ('Downtown', 'Starbucks')
```

Figure 33: Output showing Downtown station having the highest number of Starbucks outlets

How many Starbucks are there around Downtown station? It turned out that there are 4 Starbucks around this station alone.

```
Out[48]:
```

	Station	Station_Latitude	Station_Longitude	Venue	Venue_Latitude	Venue_Longitude
4115	Downtown	1.27949	103.852802	Starbucks	1.279335	103.854128
4147	Downtown	1.27949	103.852802	Starbucks	1.277949	103.850985
4166	Downtown	1.27949	103.852802	Starbucks	1.276988	103.852458
4180	Downtown	1.27949	103.852802	Starbucks	1.279422	103.854494

To find out how close these different Starbucks outlets are, I use Folium map to visualise their locations as follows.

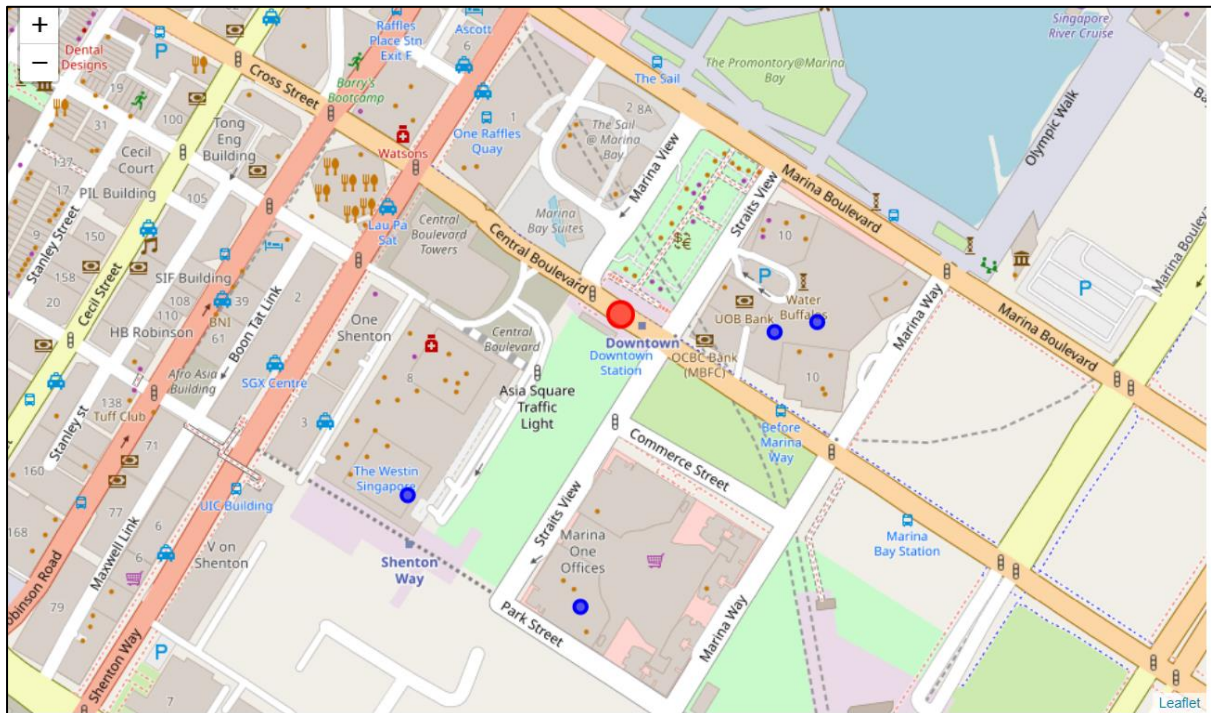


Figure 34: Map showing location of the four Starbucks outlets around Downtown station.

3.9 ATMs Around MRT Stations

Lastly, I decided to determine which stations had automatic teller machines in their vicinity (ATM) by filtering by 'Venue Category' == 'ATM'. Surprisingly, the output returned only two stations as shown below.

```
Out[50]:
```

	Station	Station_Latitude	Station_Longitude	Venue_Category	Venue	Venue_Latitude	Venue_Longitude
1132	Pioneer	1.337645	103.697420	ATM	POSB ATM	1.340155	103.697610
3959	Beauty World	1.341607	103.775682	ATM	POSB ATM @ Toh Yi	1.339966	103.773484

Figure 35: ATMs around MRT stations

4. Results

A total of 322 unique venue categories and 3223 unique venues were found around the 119 MRT stations. Analysis of the top 5 venue categories around each station found that eateries such as restaurants, coffee shops and cafes, featured prominently, proving that Singapore is indeed a foodie nation.

The clustering results were quite complex. Cluster 0 appeared to be composed of the central town regions (Raffles place, Marina, Outram Park and Orchard areas), the stations along the downtown line leading to Bukit Panjang and the stations near the ends of the East-West line, North-South line and North-East Line. Examination of the top 5 venues for this cluster revealed that they are characterised by having a Japanese restaurant. Cluster 1 consisted of just one station – Tuas West Road. This is likely to be the result of this station being a new station in an industrial area, with few venues around it. Indeed, the only venue in the top 5 listing was ‘tourist information center’ thus making it an outlier and the only member of its cluster. Cluster 2 was concentrated on suburban stations around Pasir Panjang, Bukit Batok and interestingly followed the new section of the Downtown line from Geylang Bahru all the way to Tampines East. These are often smaller housing estates. Examination of the top 5 venue categories in this cluster revealed the common presence of coffee shop and food courts. Cluster 3 consisted of older estates (Redhill, Queenstown, Rochor, Beauty World e.t.c) and some older stations along the Eastern side of the East-West line. They are characterised by having noodle houses, and an Indian/ Asian restaurant. Cluster 4 consisted only of 3 stations – Lorong Chuan, Bartley and Upper Changi. These have a bus stop and café as their top 5 venue categories.

Jalan Besar had the maximum number (21) of one venue category, i.e. Indian restaurants. Overall, the top 10 venue categories around the MRT stations are coffee shops, Chinese restaurants, café, food court, Japanese restaurant, hotel, Asian restaurant, bakery, Indian restaurant and fast food restaurant in that order. An overwhelming 9 out of the 10 top venue categories were associated with food.

The most common venue in any MRT station turned out to be the American coffee chain Starbucks, which had 4 outlets around Downtown station alone. Only two ATMs were found at Beauty World and Pioneer stations, an abnormality which would be discussed in the following section.

5. Discussion

The number of unique venue categories (322) was surprisingly large, for a small city state like Singapore. It could be that some of the venues were classified as different categories by different users (for example, Toast box could be classified as a breakfast place or a café. The clustering results for some clusters (0 and 2 in particular) appeared to follow certain sections of the MRT line, which would warrant further investigation. It appeared that these sections of the MRT network, being built around the same timeframe or surrounded by housing estates of a certain age, had similar characteristics or amenities, such as Japanese restaurants (Cluster 0) or food courts (Cluster 2).

It was surprising that Jalan Besar had 21 Indian restaurants, since it is located away from the Little India district, showing how the siting of restaurants may evolve organically, and differ from common expectation.

Workers in the downtown station area could be of an expatriate profile, since there are four Starbucks in that locale alone. A surprising number justifiable only by having sufficient numbers of coffee drinkers.

The unexpectedly small number of ATMs returned could be the result of several factors. First, ATM being a ubiquitous amenity, serves a rather mundane, though necessary function. Hence, Foursquare users are unlikely to check in and write a review since it is neither cool nor hip. The two stations with ATMs checked

in could be because it is difficult to find one in that station, hence users are motivated to leave a tip for future visitors. Another reason could be that ATMs located in malls around MRTs are so common that people would not think of checking in at that location on Foursquare.

6. Conclusion

The above analysis shows both the power and limitations of using Foursquare API location data to analyse venues around a location. While it can lead to interesting insights, as the correlation between certain sections of the MRT line with different clusters and the surprising observation of Indian restaurants concentration around Jalan Besar; for utility-type venues such as ATMs, it is not as ideal. This is due to the nature of Foursquare data, which is dependent on user check-ins. User check-ins are biased towards restaurants and hip cafes, that would be cool to share with their friends, not everyday venues such as ATMs, post offices or banks. For the latter type of venues, it would be better to use location data from the bank or post offices directly, since that would capture all the available locations for that facility.